

STUDY GUIDE · COMPLETE

Claude Certified Associate Foundations (CCAO)



Antonio Miglio

Claude Certified Associate · Foundations

[linkedin.com/in/antonio-miglio](https://www.linkedin.com/in/antonio-miglio)

PART 1

Foundations: the key concepts by domain

The 7 domains of the CCAO-F blueprint, summarized in the concepts underlying every exam question. They are not a substitute for hands-on practice with Claude, but they establish the right mental model before tackling the Compact Guide in Part 2 and the exam questions in Part 3.

Introduction — overview of the domains

The CCAO-F blueprint weighs each question according to the domain it belongs to. Output Evaluation and Validation alone accounts for a fifth of the exam: judging whether an output is good enough to use is the exam's central skill, more than producing the output in the first place.

#	CCAO-F blueprint domain	Exam weight
1	Prompting and Task Execution	14.0%
2	Output Evaluation and Validation	21.0%
3	Product and Model Selection	12.0%
4	Workflow Integration and Solution Design	16.0%
5	Configuration and Knowledge Management	12.0%
6	Governance, Risk, and Responsible Use	15.0%
7	Troubleshooting and Optimisation	10.0%

A note on the format: the real CCAO-F exam consists of 60 items in 120 minutes, a mix of multiple-choice (select one) and multiple-response (select a stated number) questions — every item states up front how many answers to select. This guide's practice set (Part 3) includes 75 — the 60 core items in the same domain proportions, plus 15 bonus items on additional topics.

Domain 1 — Prompting and Task Execution (14.0%)

The core distinction of the domain: a vague, single-shot request produces a vague, single-shot answer. Getting useful work out of Claude is a skill of framing the task, not a property of the model alone.

- **Task decomposition:** a large, ambiguous request ("improve our onboarding") is broken into smaller, well-defined sub-tasks (audit the current flow, draft new copy, propose metrics) that Claude can address one at a time, with the person reviewing between steps.
- **Iteration over one-shot:** the first draft is a starting point, not the deliverable. Refining a prompt with more context, a clearer audience, or a corrected assumption reliably beats abandoning the approach and starting over.
- **Matching strategy to task type:** analysis and research call for showing sources and asking for reasoning; drafting calls for tone and audience guidance; brainstorming calls for volume and divergence over precision. The same prompting habits do not serve every task type equally.
- **Providing context, not just instructions:** who the output is for, what has already been tried, and what "good" looks like all belong in the prompt. An instruction without context forces Claude to guess at constraints the person already knows.

- **Examples over description:** showing one or two examples of the desired output format or tone is more reliable than describing the format in the abstract, especially for structured or stylistically specific deliverables.
- **Scoping the ask to the task:** a quick internal note and a board-ready memo warrant different levels of prompting effort; over-specifying a low-stakes task wastes time, and under-specifying a high-stakes one wastes the output.

Domain 2 — Output Evaluation and Validation (21.0%)

The core distinction of the domain: fluent and confident is not the same as correct and complete. The heaviest domain on the exam is entirely about closing that gap before an output goes anywhere.

- **Hallucination awareness:** Claude can state incorrect facts, invented citations, or fabricated numbers in a tone indistinguishable from its accurate statements. A confident tone is never itself evidence of accuracy.
- **Fact-checking against a source:** specific claims, statistics, names, and dates that matter to the decision should be checked against an authoritative source before use, not accepted because they sound plausible or match expectations.
- **Completeness checking:** an output can be accurate in everything it says while silently leaving out a requirement, a stakeholder, or a risk that was in the original brief. Checking against the original request, not just against the output's own internal consistency, catches this.
- **Bias and one-sidedness:** a summary or recommendation can over-represent one perspective, one data source, or one demographic without any single sentence being factually false. Reviewing for what is missing or over-weighted is a distinct check from fact-checking.
- **Escalation thresholds:** high-stakes, irreversible, legally sensitive, or externally visible outputs warrant a second human reviewer regardless of how confident the output reads; low-stakes internal drafts do not need the same bar.
- **Editing for the audience:** an output correct for an internal technical audience may need restructuring, simplifying, or re-tuning in tone before it is appropriate for an executive, a customer, or the general public — accuracy and audience-fit are separate checks.
- **Choosing the right output format:** a structured table suits comparison, an artifact suits something the person will keep editing, and inline chat suits a quick, disposable answer — matching format to how the output will actually be used is part of validation, not an afterthought.
- **Comparing multiple outputs:** when several drafts or several models' outputs are available, comparing them against the same explicit criteria (accuracy, completeness, tone, length) produces a more defensible choice than picking on gut feel.
- **Verifying claims Claude flags as uncertain:** Claude may hedge or flag low confidence on a specific claim; treating that flag as a prompt to verify, rather than smoothing it over in the final edit, preserves the very signal that was meant to protect the reader.

Validation checklist: before you trust an output

Domain 2 alone is a fifth of the exam, and it rewards a repeatable habit more than a long list of things to remember. Before an output leaves your hands — is sent, filed, or acted on — run it through the same few checks, in the same order, every time.

- **Accuracy:** are the specific facts, figures, names, and claims verifiable against a source you trust — not merely plausible or consistent with what you expected to see?
- **Completeness:** does the output address everything the original request asked for, including requirements that were easy to overlook, not just the parts that were easiest to answer?

- **Bias and balance:** does the output fairly represent the relevant perspectives and data, or does it lean on the most available source, the loudest viewpoint, or the framing the prompt happened to suggest?
- **Internal consistency:** do the output's own numbers, dates, and claims agree with each other from start to finish, with no contradiction introduced partway through a longer piece?
- **Audience and format fit:** is this the right level of detail, tone, and structure for who will actually receive it, or was it written for a different audience by default?
- **Escalation:** given what this output will be used for — its visibility, reversibility, and stakes — does it need a second human reviewer before it goes any further, regardless of how confident it reads?

Domain 3 — Product and Model Selection (12.0%)

The core distinction of the domain: the right tool for a one-off question is rarely the right tool for a recurring workflow, and the most capable model is rarely the right default.

- **Chat vs Projects:** a single, self-contained question belongs in a plain chat; a recurring body of work with its own context, files, and conventions (a client account, a course, a product line) belongs in a Project, where that context persists across conversations.
- **Research mode for open-ended questions:** a question that needs synthesizing several sources rather than a single quick answer is a better fit for a research-oriented workflow than a single-turn chat message.
- **Artifacts for iterated deliverables:** a document, table, or piece of content the person will revise over several turns belongs in an artifact, where it can be edited in place, rather than being retyped in full in every reply.
- **Model tiers as a cost/speed/quality tradeoff:** a lighter, faster model fits high-volume, simple, well-defined tasks; a mid-tier model fits everyday drafting and analysis; the most capable tier fits complex reasoning, ambiguous requests, or high-stakes output — matching tier to task avoids both overpaying and under-delivering.
- **Context window as a finite, shared resource:** the window holds both what has been provided and what will be generated; a long-running conversation eventually needs restarting, summarizing, or moving to a Project with persistent knowledge rather than being extended indefinitely.
- **When to restart a conversation:** once a conversation has drifted across several unrelated topics, or has grown long enough that early context is unlikely to still be influencing the current answer usefully, starting fresh produces a more focused result than continuing to add to it.
- **Aligning selection with the actual constraint:** the binding constraint on a given task is sometimes cost, sometimes latency, sometimes output quality — the right product or model choice follows from identifying which constraint actually matters for that task, not from defaulting to the most powerful option available.

Domain 4 — Workflow Integration and Solution Design (16.0%)

The core distinction of the domain: using Claude well on a single task is different from designing where Claude fits inside a recurring process shared by a team.

- **Requirements analysis with Claude:** Claude can help surface and organize requirements from stakeholder notes or a messy brief, but the person remains responsible for confirming those requirements are complete and correctly prioritized before design work proceeds.
- **Research and planning support:** synthesizing background material, drafting a project plan, or comparing options are tasks Claude can accelerate substantially, provided the outputs are still checked against domain 2's evaluation standards before being acted on.

- **Process optimization:** an existing manual process (routing requests, drafting responses, compiling a recurring report) is a candidate for Claude assistance when its steps are describable and its inputs are consistent — not every process is a good fit, particularly ones requiring undocumented tacit judgment.
- **Redesigning vs bolting on:** a workflow already broken by an unclear division of responsibility is not fixed by adding Claude at one step; effective integration usually means rethinking the step's place in the workflow, not just automating it as-is.
- **Supporting solution design and iteration:** Claude can propose a first-pass solution structure and revise it against feedback, functioning as a thinking partner across iterations rather than only a single-shot generator.
- **Communicating value and limitations to stakeholders:** a stakeholder proposal for adopting Claude in a workflow should state plainly what it will and will not do, and what human oversight remains in place — overselling capability erodes trust the first time the limitation surfaces in production.
- **Piloting before scaling:** a small, reversible pilot on a real but bounded slice of the workflow reveals failure modes that a plan built purely on paper does not — scaling a workflow integration before piloting it raises the cost of a wrong assumption.
- **Ownership after handoff:** once Claude is integrated into a live workflow, someone must own monitoring its output quality over time; a successful pilot with no assigned ongoing owner tends to degrade unnoticed.

Domain 5 — Configuration and Knowledge Management (12.0%)

The core distinction of the domain: what Claude knows in a given conversation is a function of what has been explicitly configured or provided, not a property of the model in general.

- **Project instructions:** standing guidance that should apply to every conversation within a Project — tone, scope, house style, what to always check — belongs in the Project's instructions, not repeated by hand in every new chat.
- **Knowledge sources:** uploading the documents, style guides, or reference material a Project's work depends on gives Claude a stable, shared basis to draw on, instead of relying on what a person happens to paste into an individual conversation.
- **Connectors:** linking sources such as Google Drive or Gmail lets Claude reference live, current material directly, which is more reliable for information that changes than a document uploaded once and left to go stale.
- **System-level instructions as durable policy:** instructions set at the system or Project level persist and apply broadly; instructions typed into a single message apply only to that message — the two are not interchangeable, and standing rules placed only in a one-off message will not carry forward.
- **Keeping configuration current:** knowledge sources and instructions drift out of date as a team's process, product, or policy changes; a Project's configuration needs an owner who updates it, not a one-time setup treated as permanent.
- **Scoping knowledge appropriately:** a Project should hold the knowledge relevant to its own scope of work; mixing unrelated bodies of knowledge into one Project makes it harder for Claude to apply the right material and easier for irrelevant content to leak into an answer.

Domain 6 — Governance, Risk, and Responsible Use (15.0%)

The core distinction of the domain: whether a use case is technically possible with Claude is a separate question from whether it is appropriate, permitted, and safe to run.

- **Appropriate vs inappropriate use cases:** a use case can be technically straightforward and still be a poor fit for AI assistance — for example, a decision that legally or ethically requires a specific accountable human, or one resting on information too sensitive to place in a prompt at all.

- **Data sensitivity:** personal data, health information, financial details, and other regulated categories carry handling obligations that do not disappear because the processing step involves Claude rather than a person; what would need protecting on paper needs protecting in a prompt.
- **Regulatory and privacy considerations:** sector-specific rules (data protection law, sector regulation, contractual confidentiality obligations) constrain what may be shared with any AI system, including Claude, and those constraints are the organization's to know and apply, not Claude's to enforce.
- **Organizational AI policy:** an organization's own acceptable-use policy for AI tools — what data may be used, which use cases require approval, which are prohibited outright — governs individual usage decisions even where the technology itself would allow more.
- **Human accountability:** the human who acts on an AI-assisted output remains accountable for that action; "Claude produced it" does not transfer responsibility for a decision's correctness or its consequences.
- **Ethical implications beyond compliance:** an output can be fully compliant with policy and law and still raise a fairness, transparency, or dual-use concern worth weighing on its own terms — governance is broader than a checklist of rules already written down.
- **Transparency about AI involvement:** where an audience would reasonably expect to know that content was AI-assisted (a public communication, a decision affecting someone materially), disclosure is a distinct consideration from the content's accuracy.

Domain 7 — Troubleshooting and Optimisation (10.0%)

The core distinction of the domain: an underperforming result is a signal to diagnose the cause, not an invitation to immediately rewrite the prompt or discard the approach.

- **Diagnosing before rewriting:** an off-target output can stem from missing context, an ambiguous instruction, a mismatched model tier, or a genuinely hard task — identifying which one it is determines whether the fix is more context, a clearer ask, a different tier, or accepting the task needs a different approach entirely.
- **Isolating one variable at a time:** when adjusting a prompt to fix a problem, changing several things at once (wording, examples, format, model) makes it impossible to tell which change actually helped, and can hide a new problem introduced alongside the fix.
- **Using feedback to adjust approach:** a pattern across several outputs — consistently too long, consistently missing a section, consistently the wrong tone — points to a standing instruction that needs adding, not a one-off correction repeated by hand every time.
- **Recognizing task-model mismatch:** a persistently weak result on a task requiring deep reasoning or nuanced judgment may indicate the task has outgrown the model tier in use, rather than indicating the prompt itself is at fault.
- **Optimising for efficiency without losing quality:** shortening a prompt, batching similar requests, or reusing a proven instruction set can reduce time and effort on a recurring task, provided the output is re-validated after the change rather than assumed equivalent.
- **Knowing when the task, not the tool, is the problem:** some tasks are a poor fit for AI assistance regardless of prompting skill — recognizing that boundary is itself part of troubleshooting, rather than continuing to iterate on a task that needs a different kind of resource.

PART 2

Compact Guide — 24 principles to remember

Theory only: a quick review, without item-by-item detail. Read this after Part 1, before attempting the mock exam in Part 3.

1. Decomposition before delegation

A large, ambiguous request produces a large, ambiguous answer. Breaking it into named sub-tasks — each with its own clear deliverable — lets Claude address each part well and lets the person review between steps instead of only at the end.

Principle: *if a request can't be stated in one clear sentence, it should be split before it's sent, not after the output disappoints.*

2. Iteration beats restart

A first draft that misses the mark is usually one clarification away from being right — a missing constraint, an unstated audience, an example of the desired format. Refining beats discarding the approach and starting over from a blank prompt.

Principle: *diagnose what the prompt was missing before concluding the approach itself failed.*

3. Prompting effort should match task stakes

A quick internal note and a board-ready memo do not deserve the same amount of prompt engineering. Over-specifying a low-stakes task wastes time; under-specifying a high-stakes one wastes the output.

Principle: *scale the care put into a prompt to what the output will be used for, not to a fixed personal habit.*

4. Confidence is not correctness

Claude's tone does not change between an accurate statement and a fabricated one. A fluent, assured answer carries no more evidence of truth than a hesitant one — the fluency is a property of language generation, not of the underlying facts.

Principle: *never let how an answer reads substitute for checking whether it's true.*

5. Accuracy and completeness are separate checks

An output can be fully accurate in everything it states while silently leaving out a requirement, a stakeholder, or a risk from the original brief. Checking for missing content is a distinct pass from checking whether the present content is correct.

Principle: *validate against the original request twice — once for what's wrong, once for what's missing.*

6. Bias can survive a fact-check

A summary can be composed entirely of true statements and still over-represent one source, one perspective, or one demographic. Fact-checking each sentence does not catch an imbalance in which sentences were included.

Principle: *review for what the output emphasizes and omits, not only for what it asserts.*

7. Escalation follows stakes, not confidence

Whether a second human reviewer is needed depends on how visible, reversible, and consequential the output is — not on how certain the output sounds or how certain the person feels after a quick read.

Principle: *set the escalation bar from the decision's stakes, before reading the output at all.*

8. Format is a validation criterion, not a finishing touch

A structurally sound, accurate answer delivered in the wrong format — inline text where a table was needed, a static document where an editable artifact was needed — still fails the person who has to use it.

Principle: *check output format against how the result will actually be used, before checking anything else.*

9. Chat is disposable, Projects are durable

A one-off question with no lasting context belongs in a plain chat. A recurring body of work — a client account, a course, a product line — belongs in a Project, where instructions and knowledge sources persist across every future conversation about it.

Principle: *if the same context would need retyping next week, it belongs in a Project, not a chat message.*

10. Model tier follows the task's actual constraint

A lighter, faster tier suits high-volume, well-defined work; the most capable tier suits ambiguous, high-stakes, or deeply analytical work. The binding constraint — cost, speed, or quality — determines which tier is correct for a given task.

Principle: *identify which constraint actually matters for this task before choosing a tier, rather than defaulting to the most powerful option available.*

11. Context windows are a shared, finite budget

Everything already in a conversation competes for the same space as what still needs to be generated. A conversation that has drifted across unrelated topics, or grown very long, benefits more from a fresh start or a summary than from being extended further.

Principle: *a long conversation is a reason to restart or summarize, not a reason to keep adding to it indefinitely.*

12. Requirements work still needs a human owner

Claude can organize a messy brief into structured requirements, but confirming those requirements are complete, correctly prioritized, and actually agreed by stakeholders remains a human responsibility that the drafting step does not discharge.

Principle: *Claude accelerates requirements analysis; it does not replace requirements sign-off.*

13. Integrating Claude means redesigning the step, not just automating it

A workflow already broken by unclear ownership or unclear inputs does not improve by inserting Claude at the broken step. Effective integration usually means rethinking that step's place in the process first.

Principle: *fix the process design before automating the process, or the automation will faithfully reproduce the same flaw at higher speed.*

14. Pilot before you scale

A small, reversible pilot on a real but bounded slice of a workflow surfaces failure modes a plan built only on paper never will. Scaling an integration before piloting it multiplies the cost of any wrong assumption.

Principle: *never scale a workflow change past the point where a wrong assumption would still be cheap to reverse.*

15. A pilot needs an owner after it ends

A successful pilot with no one assigned to monitor its ongoing output quality tends to degrade unnoticed as inputs, staff, or requirements shift underneath it.

Principle: *integration isn't finished at launch — someone must own its quality over time, by name.*

16. Standing instructions belong at the Project level

Guidance meant to apply every time — tone, scope, a house style, a rule to always check — belongs in Project or system-level instructions, which persist. The same guidance typed into a single message applies to that message only.

Principle: *if a rule should hold for every future conversation, put it where it will still be read next time — not in today's message.*

17. Knowledge sources need an owner too

Uploaded documents and connected sources drift out of date as the underlying process, product, or policy changes. A Project's configuration is only as reliable as the last time someone updated it.

Principle: *treat a Project's knowledge sources as living material to maintain, not a one-time setup to file away.*

18. Connectors beat stale uploads for live information

A document uploaded once captures a single moment; a connector to a live source reflects what is current whenever it's queried. For information that changes, that difference matters.

Principle: *for material that changes over time, connect to the live source rather than re-uploading a snapshot.*

19. Capability and appropriateness are different questions

That Claude can technically produce an output says nothing about whether doing so is permitted, safe, or wise for this particular use case. Some decisions require a specific accountable human by policy or by law, regardless of what a model could draft.

Principle: *before asking whether Claude can do something, confirm whether it should be asked to.*

20. Sensitive data rules don't relax for AI

Personal, health, financial, and other regulated data carry the same handling obligations in a prompt that they carry anywhere else. The presence of an AI system in the workflow does not create an exception to existing data protection duties.

Principle: *if a piece of information would need protecting on paper, it needs the same protection in a prompt.*

21. Using Claude doesn't transfer accountability

The human who acts on an AI-assisted output remains responsible for that action's correctness and consequences. Attributing an error to the tool that helped draft it does not change who is accountable for having sent it.

Principle: *accountability for an action stays with the person who takes it, whatever assisted in drafting it.*

22. Compliant is not the same as ethical

An output can satisfy every written policy and legal requirement and still raise a fairness, transparency, or dual-use concern worth weighing on its own terms, separately from the compliance checklist.

Principle: *clearing policy is the floor for responsible use, not the ceiling.*

23. Diagnose before you rewrite

An off-target output can trace to missing context, an ambiguous instruction, the wrong model tier, or a task that is genuinely difficult. Rewriting the prompt at random treats all four causes the same and often fixes none of them.

Principle: *identify which of the usual causes is actually responsible before changing anything.*

24. Change one variable at a time

Adjusting wording, examples, format, and model tier all at once makes it impossible to tell which change helped, and can hide a new problem the change itself introduced.

Principle: *when troubleshooting a weak result, change a single variable and re-check before changing another.*

PART 3

Exam Questions — Mock Exam

75 original questions (60 from the core set, distributed by domain in the same proportions as the official blueprint, + 15 bonus questions on topics a standard practice set tends to leave out). Annotated solutions are in Part 4.

Important — note on NDA and originality

These are original practice questions, written from the exam's public guide and its blueprint objectives. They are not real exam content, they do not come from any live item bank, and they do not reproduce any question encountered during the exam. The content of the Claude exam is confidential and covered by an NDA, and yours will be too.

Domain 1 — Prompting and Task Execution (14.0%)

8 questions in this domain.

Question 1.1 *Multiple choice · select ONE*

A marketing coordinator at a mid-size retailer asks Claude, in one message, to "fix our email marketing." The response comes back generic and unusable. What is the most likely reason, and the best fix?

- A. Claude cannot help with marketing tasks; a specialist tool is required
- B. The request needs to be broken into specific sub-tasks — such as auditing recent open rates, drafting new subject lines, and proposing a send-time test — addressed one at a time
- C. The coordinator should repeat the exact same message a second time
- D. Switching to a more capable model tier will resolve the vagueness automatically

Question 1.2 *Multiple choice · select ONE*

A program officer at a nonprofit drafts a grant proposal with Claude. The first draft undersells the program's impact. Which response best follows sound prompting practice?

- A. Discard the draft and try an entirely different AI tool
- B. Refine the prompt with the missing information — specific outcomes, numbers, and the funder's stated priorities — and ask for a revision
- C. Accept the draft as final since it was AI-generated and reflects the available information
- D. Ask Claude to simply make the tone more impressive without adding new information

Question 1.3 *Multiple choice · select ONE*

A university department head wants to (1) generate as many possible themes for a student recruitment campaign and (2) verify the accuracy of enrollment statistics cited in a report. Which pairing of prompting strategies fits these two tasks?

- A. Use the same precise, narrowly-scoped prompt for both tasks
- B. For (1), prompt for volume and divergence; for (2), prompt for sourcing and verification, not creative range
- C. For (1), ask for a single best theme; for (2), ask for the most creative interpretation of the statistics
- D. Neither task benefits from any difference in prompting approach

Question 1.4 *Multiple choice · select ONE*

An operations manager asks Claude to "write a summary of this quarter's logistics report" and pastes the report, with no other detail. The resulting summary is technically accurate but far too detailed for its intended reader, the company's board. What was missing from the prompt?

- A. The report itself, which was already provided
- B. Context about the audience and purpose — that this is a board-level summary requiring a different level of detail than the source report
- C. A request for the summary to be shorter than the original in a general sense
- D. Nothing; the manager should simply delete unnecessary sentences after generation

Question 1.5 *Multiple choice · select ONE*

A communications team wants Claude to draft social posts that consistently match a specific, distinctive brand voice. Written descriptions of the voice haven't produced reliable results. What is the most effective next step?

- A. Provide two or three real examples of posts that capture the desired voice, alongside the request
- B. Make the written description of the voice even longer and more detailed
- C. Lower expectations, since brand voice cannot be captured in a prompt
- D. Ask a different question entirely, unrelated to brand voice

Question 1.6 *Multiple choice · select ONE*

A team lead spends significant time crafting a detailed, carefully-worded prompt for routine internal meeting notes that only three colleagues will skim once. What does this suggest about prompting practice?

- A. This is good practice; every prompt should be maximally detailed
- B. The prompting effort is likely mismatched to the task's low stakes — a lighter, faster prompt would serve just as well
- C. Meeting notes should never be drafted with AI assistance
- D. The time spent guarantees a better outcome regardless of the task

Question 1.7 *Multiple choice · select ONE*

An events coordinator is planning a 300-person conference and asks Claude, in a single request, to "handle the event planning." Which approach best applies task decomposition?

- A. Repeat the same broad request with more enthusiastic wording
- B. Split the request into named sub-tasks — venue shortlist, draft agenda, vendor outreach email, budget tracker — and address each with its own prompt and review
- C. Ask Claude to pick which sub-tasks matter most and only do those, without coordinator input
- D. Wait until closer to the event date to begin any prompting

Question 1.8 *Multiple response · select TWO*

Which TWO of the following reflect sound practice for matching a prompting strategy to the type of task at hand?

- A. Request the widest possible range of options for an open brainstorming task, rather than a single early favorite
- B. Use the exact same prompt structure for fact-verification tasks and creative brainstorming tasks
- C. For a drafting task, specify the intended audience and tone rather than leaving them unstated
- D. Assume every task benefits equally from the same level of prompting detail

- E. Treat a request for options and a request for a verified fact as needing identical handling

Domain 2 — Output Evaluation and Validation (21.0%)

13 questions in this domain.

Question 2.1 *Multiple choice · select ONE*

A research assistant asks Claude to summarize academic literature on a topic and include citations. The summary reads fluently and the citations look properly formatted, but two of them turn out not to exist when checked. What does this illustrate?

- A. Properly formatted citations are inherently trustworthy
- B. Claude can generate fabricated content, including citations, in a tone indistinguishable from accurate content — formatting confidence is not evidence of existence
- C. This can only happen with obscure or highly technical topics
- D. The assistant made an error in copying the citations, not Claude

Question 2.2 *Multiple choice · select ONE*

A finance-adjacent operations report drafted with Claude states a specific industry growth statistic that will be quoted in an external presentation. What should happen before that number is used?

- A. Nothing further, since the number reads naturally within the report
- B. The statistic should be checked against an authoritative source before it is used externally
- C. The number should be rounded to reduce the appearance of false precision
- D. The statistic should be removed entirely rather than checked, out of caution

Question 2.3 *Multiple choice · select ONE*

A project manager asked Claude to summarize stakeholder feedback from six interviews. The resulting summary is accurate in every sentence but never mentions the one interviewee who raised a budget concern. What kind of failure is this?

- A. A hallucination, since something was included that shouldn't have been
- B. A completeness failure — accurate content that silently omits material that was part of the original input
- C. Not a failure at all, since every included sentence is factually correct
- D. A formatting failure that has nothing to do with the summary's content

Question 2.4 *Multiple choice · select ONE*

A summary of customer feedback draws almost entirely on the most vocal, most negative reviews, even though the majority of feedback collected was neutral to positive. Every quote and statistic in the summary is accurate. What is the concern here?

- A. There is no concern, since nothing false was stated
- B. The summary may be biased in its selection and emphasis, even without containing any factual error
- C. The summary should have included the star ratings for extra credibility
- D. This is purely a formatting issue

Question 2.5 *Multiple choice · select ONE*

A Claude-drafted response to a regulatory inquiry reads clearly and confidently, and the drafting employee feels satisfied with it. Should it be sent without further review?

- A. Yes, since the drafter's confidence in the output is itself a sufficient check
 - B. No — a regulatory response is high-stakes and externally visible, which calls for a second human reviewer regardless of how confident the draft reads
 - C. Yes, but only if the draft is under one page long
 - D. No, but only because AI-assisted drafts must never be sent externally under any circumstances
-

Question 2.6 *Multiple choice · select ONE*

An analyst produces a technically accurate summary of a complex data migration for an internal engineering audience, then needs to give the same information to the company's executive team. Is the same document appropriate for both audiences?

- A. Yes, since the underlying facts are identical
 - B. Not necessarily — accuracy and audience-fit are separate checks, and the technical version may need restructuring or simplifying for executives
 - C. No, an entirely new set of facts must be generated for executives
 - D. Yes, executives should be expected to read technical material as delivered
-

Question 2.7 *Multiple choice · select ONE*

A team needs to compare pricing across twelve vendors, and separately needs a one-off answer to a quick factual question. Which choices of output format best fit these two needs?

- A. A structured table for the comparison; a quick inline chat answer for the one-off question
 - B. An artifact document for both, regardless of the difference in ongoing use
 - C. Inline chat text for both, since format does not affect usability
 - D. A structured table for the one-off question and inline chat for the vendor comparison
-

Question 2.8 *Multiple choice · select ONE*

Two drafts of a client-facing announcement are available, generated with different phrasing. A reviewer needs to choose one. What is the most defensible way to decide?

- A. Pick whichever one reads more fluently on a first skim
 - B. Compare both drafts against the same explicit criteria — accuracy, completeness, tone, and length — and choose based on that comparison
 - C. Ask a colleague to flip a coin
 - D. Combine both drafts into one without re-reading either for consistency
-

Question 2.9 *Multiple choice · select ONE*

While drafting a report, Claude flags one specific figure as uncertain and suggests double-checking it. In the final edit, the employee smooths over the hedge so the report reads more confidently, without actually verifying the figure. What went wrong?

- A. Nothing — a more confident tone always improves a final report
 - B. The hedge was a genuine signal to verify that figure; removing it without checking discarded the very protection it was meant to provide
 - C. The employee should have asked Claude to remove the hedge earlier in the process
 - D. Flagged uncertainty from Claude is never worth acting on
-

Question 2.10 *Multiple response · select TWO*

A vendor-selection memo is accurate in every stated fact. Which TWO additional checks are still needed before it can be considered validated?

- A. Confirming the memo addresses every criterion the original request specified, not only the ones easiest to answer
- B. Re-reading it once more for typos
- C. Reviewing whether the memo fairly represents all the vendors considered, rather than favoring the one the drafter expected to win
- D. Making the memo longer to appear more thorough
- E. Nothing further, since accuracy is the only thing that matters

Question 2.11 *Multiple choice · select ONE*

A lengthy financial summary states a total of \$4.2 million in one section and a total of \$4.6 million for the same category two pages later. Every individual figure, checked separately, traces back correctly to source data. What kind of problem is this?

- A. No problem, since each number is individually sourced correctly
- B. An internal consistency failure — the document contradicts itself even though each piece traces to a real source
- C. A formatting issue unrelated to content
- D. A bias issue, since one number is being favored over the other

Question 2.12 *Multiple response · select TWO*

A summary of employee survey results reads fluently, uses correctly-quoted figures, and was generated quickly. Which TWO of the following would still need to be checked before distributing it company-wide?

- A. Whether it fairly represents both positive and negative feedback in proportion to what was actually collected
- B. Whether every sentence uses varied vocabulary
- C. Whether the level of detail and tone suit a company-wide audience, or whether it was written more like an internal analyst note
- D. Whether the document is under a specific word count
- E. Nothing, since correctly-quoted figures are sufficient validation on their own

Question 2.13 *Multiple choice · select ONE*

An intern feels very confident that a Claude-drafted customer apology email is ready to send, describing it as "clearly well-written." The email will go to a customer considering legal action. What should determine whether it needs further review?

- A. The intern's confidence in the writing quality
- B. How well-written the email appears on a first read
- C. The stakes and reversibility of sending it to a customer in this specific situation — which call for review regardless of how the draft reads
- D. Whether the email is shorter than a page

Domain 3 — Product and Model Selection (12.0%)

7 questions in this domain.

Question 3.1 *Multiple choice · select ONE*

A consultant manages ongoing work for one long-term client, involving recurring documents, a specific tone, and reference material used across many conversations over several months. Which setup fits best?

- A. A fresh, unrelated chat for every new question about the client
- B. A Project dedicated to the client, holding the reference material and standing instructions so they persist across conversations
- C. Pasting the same background information into every new chat by hand
- D. None of these — recurring client work cannot be supported this way

Question 3.2 *Multiple choice · select ONE*

An analyst needs to answer "what are the current best practices for reducing customer churn in subscription businesses," which requires synthesizing many sources rather than a single quick fact. Which approach fits best?

- A. A single-turn chat message with no further elaboration
- B. A research-oriented workflow suited to synthesizing multiple sources on an open-ended question
- C. Asking the question repeatedly until a satisfying answer appears
- D. Restricting the question to a single named source only

Question 3.3 *Multiple choice · select ONE*

A team will revise a partnership proposal across many rounds of feedback over two weeks, with each round changing specific sections. What is the most efficient way to work with Claude on this?

- A. Ask Claude to retype the entire proposal in full in the chat window after every single round of feedback
- B. Use an artifact so the proposal can be edited in place across rounds, rather than regenerated from scratch each time
- C. Start an entirely new conversation for every round of edits
- D. Avoid using Claude for anything requiring more than one round of revision

Question 3.4 *Multiple choice · select ONE*

A support team needs to auto-tag thousands of incoming tickets per day into simple categories, a well-defined, high-volume, low-complexity task. Which consideration should most influence model tier selection here?

- A. Always select the most capable tier available, regardless of task complexity
- B. A lighter, faster tier is likely the better fit, since the task is simple and high-volume, and using a heavier tier would add cost and latency without improving accuracy
- C. Model tier makes no difference to cost or speed
- D. The tier should be chosen at random for each ticket to increase variety

Question 3.5 *Multiple choice · select ONE*

A conversation with Claude has run for several hours, drifted across three unrelated topics, and grown very long. The person wants to ask a focused new question on a fourth, unrelated topic. What is the best approach?

- A. Keep adding to the same long conversation indefinitely
- B. Start a fresh conversation (or move to a Project with the relevant persistent knowledge) for the new, unrelated topic
- C. Ask Claude to forget everything mentioned so far, within the same conversation, and continue there
- D. Make the question longer to compensate for the cluttered context

Question 3.6 *Multiple choice · select ONE*

A team is choosing between two approaches for a task where speed of response during a live customer interaction is the overriding concern, more than raw output sophistication. What should drive the product or model choice?

- A. Whichever option is most capable overall, regardless of latency
- B. The actual binding constraint for this task — here, response speed — rather than a general preference for the most powerful option
- C. The option that is most expensive, on the assumption that cost reflects quality
- D. A coin flip, since all options are functionally identical

Question 3.7 *Multiple response · select TWO*

Which TWO of the following reflect sound product and model selection practice?

- A. Defaulting to the most capable model tier for every task regardless of complexity, to avoid ever thinking about tradeoffs
- B. Moving a recurring body of work with persistent context into a Project rather than repeating setup in every new chat
- C. Matching model tier to the task's actual constraint — cost, speed, or quality — rather than to a fixed habit
- D. Extending a single conversation indefinitely rather than ever restarting or summarizing it
- E. Using a plain chat message for a task that will be revised across many rounds of feedback

Domain 4 — Workflow Integration and Solution Design (16.0%)

10 questions in this domain.

Question 4.1 *Multiple choice · select ONE*

A property management company gives Claude a messy set of stakeholder notes from a tenant-services redesign meeting and asks it to organize them into structured requirements. Claude produces a clean, well-organized list. What should happen next?

- A. Proceed directly to design work, since the list is well-organized
- B. A human should confirm the requirements are complete and correctly prioritized before design work proceeds
- C. Discard the structured list, since Claude cannot be trusted with requirements at all
- D. Ask Claude to also make the final prioritization decision independently

Question 4.2 *Multiple choice · select ONE*

A logistics coordinator uses Claude to synthesize background research and draft a project plan for a new warehouse rollout. The plan looks thorough and well-structured. What should happen before it is adopted?

- A. Nothing further, since a well-structured plan is sufficient on its own
- B. The plan should still be checked against the same evaluation standards as any other output — accuracy, completeness, and fit for the actual situation
- C. The plan should be adopted only if it is the longest of several drafts
- D. Research and planning tasks are exempt from the usual evaluation domain

Question 4.3 *Multiple choice · select ONE*

A hospital administration team is considering two candidate processes for Claude assistance: (1) drafting routine appointment-reminder messages from a fixed template, and (2) making individual patient triage

priority decisions based on undocumented clinical judgment developed over years. Which is the better candidate, and why?

- A. Both are equally good candidates for Claude assistance
- B. The reminder messages, because the process is describable and its inputs are consistent, unlike the triage decisions, which rest on undocumented tacit judgment
- C. The triage decisions, because they are more valuable to automate
- D. Neither is appropriate for any AI assistance whatsoever

Question 4.4 *Multiple choice · select ONE*

A retailer's customer-complaint process is already slow because it's unclear which department owns a complaint once it's filed. The company adds a Claude-drafted acknowledgment email at intake, without addressing the ownership ambiguity. What is the likely result?

- A. The underlying delay is fixed, since a faster acknowledgment message now exists
- B. The process is likely still slow overall, since the acknowledgment doesn't resolve the unclear ownership causing the underlying delay
- C. Adding any Claude assistance always resolves workflow bottlenecks
- D. The complaint volume will decrease as a direct result

Question 4.5 *Multiple choice · select ONE*

A product team asks Claude to propose a first-pass structure for a new customer onboarding flow, then gives feedback and asks for a revision, repeating this over several rounds. What role is Claude playing here?

- A. A single-shot generator producing one final, unchangeable answer
- B. A thinking partner across iterations, proposing and revising a solution structure in response to feedback
- C. An unrelated tool that cannot support solution design
- D. A replacement for the product team's own design judgment

Question 4.6 *Multiple choice · select ONE*

A department head is pitching stakeholders on using Claude to help draft first-pass responses to customer inquiries. The pitch claims the tool will "handle customer service end to end with no further oversight." What is the concern with this pitch?

- A. There is no concern; strong claims build stakeholder confidence
- B. It oversells capability and omits the human oversight that should remain in place, which will erode trust once a limitation surfaces in production
- C. The pitch should have used even stronger language to secure buy-in
- D. Stakeholder communication is not part of workflow integration

Question 4.7 *Multiple choice · select ONE*

A regional retail chain wants to roll out Claude-assisted inventory reporting across all 40 stores simultaneously, based on a plan that has not yet been tested anywhere. What is the recommended approach instead?

- A. Proceed with the full 40-store rollout as planned, since the plan looks sound on paper
- B. Run a small, reversible pilot at one or two stores first to surface failure modes before scaling further
- C. Skip piloting entirely to save time, since delays are costly
- D. Only roll out to the single best-performing store, permanently, with no plan to expand

Question 4.8 *Multiple choice · select ONE*

A successful three-month pilot integrating Claude into a company's expense-report review process ends, and no one is assigned to monitor its output quality afterward. Six months later, quality has quietly declined as expense categories changed. What was missing?

- A. Nothing was missing; a successful pilot needs no further attention
- B. An assigned owner responsible for monitoring the integration's output quality on an ongoing basis after the pilot ended
- C. A longer pilot period would have prevented this on its own
- D. This kind of decline is unrelated to the pilot ending

Question 4.9 *Multiple response · select TWO*

Which TWO of the following are recommended practices when integrating Claude into an existing business process?

- A. Piloting the change on a small, reversible slice of the process before scaling company-wide
- B. Adding Claude to a broken step without addressing why the step is broken
- C. Assigning an owner to monitor output quality on an ongoing basis after launch
- D. Rolling out to the entire organization at once, without a pilot, to save time
- E. Assuming a successful pilot needs no further attention once it concludes

Question 4.10 *Multiple choice · select ONE*

A publishing company's editorial workflow is unclear about who has final sign-off on a manuscript once AI-assisted copyediting is added. Leadership wants to add Claude assistance at the copyediting step without resolving the sign-off ambiguity. What is the most likely outcome, based on sound workflow-integration practice?

- A. The ambiguity will resolve itself once Claude is added to the process
- B. The underlying ambiguity is likely to persist or worsen, since adding a capability at one step doesn't redesign the surrounding process
- C. Adding Claude always simplifies unclear ownership automatically
- D. This is not a workflow integration concern at all

Domain 5 — Configuration and Knowledge Management (12.0%)

7 questions in this domain.

Question 5.1 *Multiple choice · select ONE*

A team wants every conversation within their content-review Project to always check for a specific style rule and always avoid a specific phrase. Where should this guidance be placed?

- A. Typed into a single chat message, to be repeated by hand whenever needed
- B. In the Project's instructions, so it applies automatically to every conversation within the Project
- C. Nowhere — this kind of standing rule cannot be configured
- D. In a separate document that is never actually shared with Claude

Question 5.2 *Multiple choice · select ONE*

A legal-adjacent team's Project involves referencing a specific internal style guide and glossary of approved terminology on every task. What is the best way to make this material consistently available?

- A. Paste the entire style guide into the chat window at the start of every new conversation
 - B. Upload the style guide and glossary as knowledge sources within the Project
 - C. Rely on Claude's general knowledge, since it should already know the team's internal terminology
 - D. Email the style guide to colleagues instead of configuring the Project
-

Question 5.3 *Multiple choice · select ONE*

A sales operations team wants Claude to always reference the current version of a frequently-updated pricing sheet stored in their shared drive. Which approach is most reliable?

- A. Upload a copy of the pricing sheet once and never update it
 - B. Use a connector to the shared drive so Claude references the live, current version whenever needed
 - C. Retype the current prices into the prompt manually each time
 - D. Rely on memory of the prices from a previous conversation
-

Question 5.4 *Multiple choice · select ONE*

A manager types "always flag any mention of a competitor by name" into the middle of an ordinary chat message about an unrelated task. In the next, separate conversation, Claude does not apply this rule. Why not?

- A. This is a bug; instructions in any message should always persist to future conversations
 - B. Instructions typed into a single message apply only to that message; a durable rule needs to be set at the system or Project level to persist across conversations
 - C. Claude deliberately ignores instructions about competitors specifically
 - D. The manager needs to repeat the instruction twice in the same message for it to persist
-

Question 5.5 *Multiple choice · select ONE*

A Project's knowledge sources include a company policy document that was uploaded a year ago and has since been revised twice, with neither revision reflected in the Project. What is the risk here?

- A. None — once uploaded, a knowledge source stays accurate indefinitely
 - B. Claude may reference the outdated version of the policy, since the Project's configuration is only as current as its last update
 - C. The risk only exists if the document is explicitly opened during the conversation
 - D. Uploaded documents automatically update themselves when the source changes
-

Question 5.6 *Multiple choice · select ONE*

A single Project has been used to hold knowledge sources for three unrelated initiatives — a hiring plan, a facilities renovation, and a customer survey — because it was convenient to set up once. What problem does this create?

- A. No problem; more knowledge in one place is always better
 - B. It becomes harder for Claude to apply the right material to a given task, and irrelevant content from one initiative can leak into an answer about another
 - C. This actually improves Claude's ability to cross-reference unrelated topics
 - D. Projects have no limit on scope, so this is the recommended approach
-

Question 5.7 *Multiple response · select TWO*

Which TWO of the following reflect sound configuration and knowledge management practice?

- A. Placing standing, repeat-every-time guidance in Project instructions rather than a single chat message
- B. Uploading a knowledge source once and assuming it will always remain accurate with no further attention
- C. Using a connector rather than a static upload for information that changes frequently
- D. Combining several unrelated initiatives' knowledge sources into a single Project for convenience
- E. Typing a durable, standing rule into an individual message instead of a persistent setting

Domain 6 — Governance, Risk, and Responsible Use (15.0%)

9 questions in this domain.

Question 6.1 Multiple choice · select ONE

An HR manager considers asking Claude to draft talking points to help prepare for a difficult conversation with an employee, and separately considers asking Claude to decide whether that employee should be terminated. What is the key distinction between these two use cases?

- A. Both are equally appropriate, since Claude can produce a confident answer to either
- B. Drafting talking points is a reasonable use of assistance; the termination decision itself requires a specific accountable human and should not be delegated to Claude
- C. Neither use case is appropriate under any circumstances
- D. The termination decision is more appropriate because it saves the manager more time

Question 6.2 Multiple choice · select ONE

An operations employee is drafting a Claude prompt that would include a customer's full medical history to help summarize a case file. What should guide this decision?

- A. Medical history can be freely included, since the summary itself is for internal use only
- B. Health information carries the same handling obligations in a prompt that it carries anywhere else, and its use should be evaluated against those obligations before including it
- C. Handling rules for medical information only apply to paper records, not digital prompts
- D. There is no meaningful difference between medical information and any other type of data

Question 6.3 Multiple choice · select ONE

A company operating in multiple countries wants to use Claude to process customer data that is subject to strict regional data protection law. Who is responsible for ensuring this use complies with that law?

- A. Claude, since it is the system doing the processing
- B. The organization, which must know and apply the regulatory constraints on what can be shared with any AI system, including Claude
- C. No one, since AI processing falls outside existing data protection law entirely
- D. The individual employee typing the prompt, personally and exclusively, with no organizational responsibility

Question 6.4 Multiple choice · select ONE

An organization's AI usage policy requires manager approval before using Claude on any task involving personally identifiable customer information. An employee has a task that would technically work well with Claude but involves this kind of data, and has not sought approval. What should happen?

- A. Proceed with the task, since Claude is technically capable of handling it well
- B. Seek the required approval before proceeding, since the organization's own policy governs this decision regardless of technical capability

- C. Ignore the policy if the task is urgent enough
- D. Assume the policy does not apply because the data will only be used briefly

Question 6.5 *Multiple choice · select ONE*

An analyst sends a client a Claude-drafted figure that turns out to be incorrect, and later explains to their manager that "the AI got it wrong," implying the error isn't theirs. Is this a reasonable position?

- A. Yes, since the AI generated the incorrect content
- B. No — the person who acted on and sent the output remains accountable for its correctness and consequences, regardless of what tool assisted in drafting it
- C. Yes, but only if the analyst did not have time to check the figure
- D. This depends entirely on how the client reacts

Question 6.6 *Multiple choice · select ONE*

A marketing team's plan to use Claude for hyper-targeted ad copy aimed specifically at financially vulnerable individuals is fully compliant with current advertising law and company policy. Should that settle the matter?

- A. Yes, compliance with law and policy is the only consideration that matters
- B. Not necessarily — an output can be fully compliant and still raise a fairness or ethical concern worth weighing on its own terms
- C. No, because this use case is technically impossible for Claude to support
- D. Yes, but only because the campaign is expected to perform well

Question 6.7 *Multiple choice · select ONE*

A local government office plans to publish a set of AI-assisted answers to frequently asked questions on its public website, where residents would reasonably expect to know whether they're reading AI-generated content. What consideration applies here, separate from the content's accuracy?

- A. None — if the content is accurate, no further consideration is needed
- B. Whether to disclose that the content was AI-assisted, given that the audience would reasonably expect to know
- C. Whether the content is entertaining enough for a public audience
- D. Whether the content took less time to produce than a human-written version

Question 6.8 *Multiple response · select TWO*

Which TWO of the following are governance red flags, even if the underlying content Claude produces is accurate?

- A. Using Claude to draft a decision that, by policy, requires a specific accountable human to make
- B. Asking Claude to summarize a public news article
- C. Including regulated personal data in a prompt without considering applicable data protection obligations
- D. Using Claude to draft a first-pass outline for an internal team meeting
- E. Asking Claude to proofread a piece of publicly available marketing copy

Question 6.9 *Multiple choice · select ONE*

An employee at an insurance company wants to use Claude to help draft an internal explanation of a newly passed regulation, and separately wants to use it to auto-generate the final claims-denial decision for individual policyholders. How should these two requests be treated?

- A. Identically, since both involve Claude and the insurance industry
- B. Differently — drafting an internal explanation is a reasonable assistive use, while an individual claims-denial decision is a consequential, individually-impactful decision that calls for specific human accountability
- C. Both should be denied outright, since insurance is a regulated industry
- D. Both should be approved automatically, since regulation only concerns marketing claims

Domain 7 — Troubleshooting and Optimisation (10.0%)

6 questions in this domain.

Question 7.1 *Multiple choice · select ONE*

A weekly status report generated with Claude consistently misses the mark. Before changing anything, what is the recommended first step?

- A. Immediately rewrite the prompt with entirely different wording
- B. Diagnose the likely cause first — missing context, an ambiguous instruction, the wrong model tier, or a genuinely difficult task — before deciding what to change
- C. Switch to a different AI tool without further investigation
- D. Give up on using Claude for this report

Question 7.2 *Multiple choice · select ONE*

To fix a weak result, an employee simultaneously rewords the prompt, adds three new examples, changes the requested output format, and switches to a different model tier — all in one attempt. The new result is better. What is the problem with this approach?

- A. There is no problem; the result improved, which is all that matters
- B. It's now impossible to tell which of the four changes actually caused the improvement, and a new problem introduced by one of them could be hidden by the others
- C. This is always the fastest way to fix a prompt
- D. Changing several things at once is required for any meaningful improvement

Question 7.3 *Multiple choice · select ONE*

Over several weeks, a manager notices that Claude-drafted client emails are consistently too long, every single time, regardless of who requests them. What is the most efficient fix?

- A. Manually shorten every email by hand after it's generated, indefinitely
- B. Add a standing instruction — at the system or Project level — specifying the desired length, so it applies automatically going forward
- C. Ignore the pattern, since each email is only slightly too long
- D. Ask a different colleague to draft emails instead

Question 7.4 *Multiple choice · select ONE*

A task requiring nuanced ethical judgment across ambiguous, competing considerations consistently produces shallow, unsatisfying results even after several rounds of careful prompt refinement. What should this pattern suggest?

- A. The prompt simply needs to be worded more cleverly

- B. The task may have outgrown the model tier in use, rather than the prompt itself being at fault
- C. This kind of task is impossible for any AI assistance regardless of tier
- D. The employee should stop trying entirely

Question 7.5 *Multiple choice · select ONE*

A team shortens and reuses a proven prompt structure for a recurring weekly task to save time. What should happen after making this efficiency change?

- A. Nothing further — once a prompt is proven, no future output from it needs checking
- B. The output should still be re-validated after the change, rather than assumed automatically equivalent to before
- C. The change should be reversed immediately regardless of the result
- D. The team should stop tracking output quality for this task altogether

Question 7.6 *Multiple response · select TWO*

Which TWO of the following are sound troubleshooting practices when a Claude-assisted output consistently underperforms?

- A. Diagnosing the likely cause — context, instruction clarity, model tier, or task difficulty — before changing anything
- B. Changing the wording, examples, and model tier all at once to save time
- C. Changing one variable at a time and re-checking the result before changing another
- D. Assuming the task must be impossible for AI assistance after a single attempt
- E. Repeating the exact same prompt again without any change, expecting a different result

Bonus — Extended Topics (15 questions)

Questions that go beyond the 7 standard domains, designed to cover topics a typical practice set tends to leave out.

Question B.1 *Multiple choice · select ONE · scope: Prompting and Task Execution*

A communications officer needs a press release drafted in English and then adapted, not just translated, for a Spanish-speaking regional audience with different cultural references. Which prompting approach fits best?

- A. Draft the English version, then run it through a literal word-for-word translation with no further guidance
- B. Draft the English version, then explicitly ask for a culturally adapted version for the Spanish-speaking audience, not a literal translation
- C. Write both versions completely independently with no connection between them
- D. Assume the English version will work equally well for any audience once translated

Question B.2 *Multiple choice · select ONE · scope: Output Evaluation and Validation*

Claude generates a bar chart summarizing quarterly sales figures from data provided in the prompt. The chart looks clean and professional. What still needs to be checked before it's used in a report?

- A. Nothing — a clean, professional appearance is sufficient validation for a chart
- B. That the chart's values, axis labels, and scale accurately reflect the underlying data provided
- C. Only whether the color scheme matches the company's branding
- D. Whether the chart is visually more attractive than the previous quarter's chart

Question B.3 *Multiple choice · select ONE · scope: Product and Model Selection*

A team needs to turn a lengthy strategy document into a set of presentation slides for a leadership meeting next week, and will need to keep revising the slides for several days beforehand. Which approach is best suited to this need?

- A. Retype the entire slide content from scratch in chat every time a change is needed
- B. Use an editable, artifact-based or presentation-oriented workflow that supports revising the same deck across multiple sessions
- C. Start a brand new, unrelated conversation for every single slide
- D. Avoid using Claude for this task since it involves visual formatting

Question B.4 *Multiple choice · select ONE · scope: Workflow Integration and Solution Design*

Marketing and product teams both need to reference the same set of customer persona documents while working on related but separate deliverables. What is a reasonable configuration choice?

- A. Each team should maintain its own separate, duplicate copy of the persona documents with no coordination
- B. A shared Project (or shared knowledge sources) that both teams can draw on, keeping the underlying material consistent across their separate work
- C. Neither team should use Claude until the persona documents are merged into a single deliverable
- D. The personas should be re-created from memory by each team independently

Question B.5 *Multiple choice · select ONE · scope: Configuration and Knowledge Management*

A customer support lead wants Claude to reference the current status of open tickets from the team's ticketing system whenever asked, rather than a snapshot from whenever a document happened to be uploaded. What best supports this?

- A. A one-time export of tickets uploaded as a static knowledge source
- B. A connector to the ticketing system, so current ticket status can be referenced as it changes
- C. Asking colleagues to manually report ticket status in the chat each time
- D. This kind of live reference isn't possible with current configuration options

Question B.6 *Multiple choice · select ONE · scope: Governance, Risk, and Responsible Use*

A freelance writer plans to publish a Claude-assisted article without disclosing AI involvement anywhere, in a context where the publication's own editorial policy requires disclosure of AI-assisted content. What should the writer do?

- A. Publish without disclosure, since the article reads well regardless of how it was produced
- B. Follow the publication's editorial policy and disclose the AI assistance as required
- C. Disclosure is never required for any AI-assisted writing under any circumstances
- D. Only disclose if a reader specifically asks about it after publication

Question B.7 *Multiple choice · select ONE · scope: Troubleshooting and Optimisation*

A back-office team manually drafts the same type of short acknowledgment letter, with only a few details changing, roughly 50 times a week. What is a sound way to make this more efficient without sacrificing quality?

- A. Draft each letter completely from scratch every time, with no reuse of prior work

- B. Develop a reliable prompt template for the letter type, reuse it across instances, and periodically re-check outputs to confirm quality holds
 - C. Stop checking any individual letter once the template has been used successfully a few times
 - D. Increase the length of every letter regardless of content, to appear more thorough
-

Question B.8 *Multiple choice · select ONE · scope: Output Evaluation and Validation*

An employee asked Claude to "tighten up" a two-paragraph client update. The returned draft is well-written but has grown to five paragraphs and now includes new claims about future plans that were never part of the original message. What should this prompt?

- A. Send it as-is, since the added content reads professionally
 - B. Recognize this as scope creep — the output has grown beyond and diverged from what was actually asked — and revise it back toward the original request before using it
 - C. Assume the new claims about future plans must be accurate since they read confidently
 - D. Ask for the draft to be made even longer to match the new content added
-

Question B.9 *Multiple choice · select ONE · scope: Governance, Risk, and Responsible Use*

A boutique agency delivers a Claude-assisted market research report to a client without mentioning that AI tools were used in its production, even though the client's contract requires disclosure of any AI involvement in deliverables. What is the concern?

- A. There is no concern, since the report's content is accurate
 - B. This breaches a contractual disclosure obligation, which is separate from whether the report's content is accurate
 - C. Disclosure requirements never apply to research reports specifically
 - D. The concern only applies if the client specifically asks after the fact
-

Question B.10 *Multiple choice · select ONE · scope: Prompting and Task Execution*

A team wants to gauge whether a Claude-drafted customer response is actually an improvement over what a human agent typically writes for the same type of inquiry. What is a reasonable way to check this?

- A. Assume it is automatically better, since it was produced faster
 - B. Compare the Claude-drafted response against a sample of human-written responses using the same explicit criteria (accuracy, tone, resolution rate)
 - C. Skip comparison entirely, since AI-generated text cannot be meaningfully benchmarked
 - D. Judge based only on which response is longer
-

Question B.11 *Multiple choice · select ONE · scope: Workflow Integration and Solution Design*

A team is deciding whether to use Claude to help make final decisions in a novel, high-stakes situation with no precedent, ambiguous facts, and significant consequences for the people involved, where organizational policy requires a senior human decision-maker for exactly this kind of situation. What is the appropriate role for Claude here?

- A. Making the final decision directly, since it can process the ambiguous facts quickly
- B. Supporting the human decision-maker by organizing the facts and surfacing considerations, while the final decision itself stays with the senior human as policy requires
- C. Refusing to be involved in any capacity, even in a supporting role
- D. Making the decision, then having a human sign off afterward without genuine review

Question B.12 *Multiple choice · select ONE · scope: Product and Model Selection*

A long Project conversation supporting a multi-month initiative has accumulated so much back-and-forth that responses have started referencing details that are no longer relevant to the current phase of work. What is a reasonable response?

- A. Keep the conversation exactly as is indefinitely, since a Project's context should never be touched
- B. Summarize or reset the conversation's active context to the current phase, while keeping the Project's durable knowledge sources and instructions intact
- C. Delete the entire Project, including its knowledge sources and instructions
- D. Switch to a completely different Project unrelated to the initiative

Question B.13 *Multiple choice · select ONE · scope: Configuration and Knowledge Management*

A Project's knowledge sources include a document with sensitive compensation data, and the Project is shared broadly with a team that includes members who should not have access to that specific document. What is the underlying issue?

- A. There is no issue, since the Project is convenient to share as a whole
- B. Sharing at the Project level can expose sensitive material to people who shouldn't have access to it, so scope and access need to be considered together, not just convenience
- C. Sensitive data should never be uploaded as a knowledge source under any circumstances, for any team
- D. This is solely a data-sensitivity issue and has nothing to do with configuration

Question B.14 *Multiple response · select TWO · scope: Governance, Risk, and Responsible Use*

A small business owner is building a public-facing FAQ chatbot with light no-code AI assistance. Which TWO considerations are part of responsible use here?

- A. Whether customers interacting with the bot would reasonably expect to know they're talking with an AI system
- B. Whether the bot's interface uses the business's brand colors
- C. Whether the bot might give financial or legal-sounding advice it isn't qualified to give, without a human escalation path
- D. Whether the bot responds within two seconds on average
- E. Whether the bot was built using a no-code tool rather than custom code

Question B.15 *Multiple response · select TWO · scope: Troubleshooting and Optimisation*

A recurring analysis task has produced consistently shallow results for weeks despite several rounds of prompt refinement, each changing only one thing at a time. Which TWO conclusions are reasonable to consider at this point?

- A. The task may call for a more capable model tier than the one in use
- B. The employee should have changed several variables simultaneously from the very first attempt
- C. The task itself may be a poor fit for AI assistance in its current form, and may need to be restructured or handled differently
- D. No further diagnosis is warranted since prompting alone is always the answer
- E. The result must be acceptable simply because several rounds of effort have already been invested

PART 4

Answers and Explanations

The reasoning behind every correct answer, with a note on why the other options don't hold up. Best consulted after attempting Part 3 with a clear head, domain by domain.

Domain 1 — Prompting and Task Execution (14.0%)

1.1 — Correct: B

✓ **B. The request needs to be broken into specific sub-tasks — such as auditing recent open rates, drafting new subject lines, and proposing a send-time test — addressed one at a time**

"Fix our email marketing" has no defined deliverable. Decomposing it into named, concrete sub-tasks gives Claude something specific to act on and gives the coordinator a checkpoint after each one.

Why not the others: *A is false — the task is well within scope once it's specific. C repeats the same underlying problem. D may improve fluency but does not supply the missing specificity.*

1.2 — Correct: B

✓ **B. Refine the prompt with the missing information — specific outcomes, numbers, and the funder's stated priorities — and ask for a revision**

An underwhelming first draft is usually missing a specific input, not evidence of a failed approach. Supplying the concrete detail that was missing and iterating produces a stronger draft than starting over.

Why not the others: *A abandons a fixable approach. C accepts a weak result without attempting an easy fix. D asks for polish without addressing the actual gap — missing substance, not missing style.*

1.3 — Correct: B

✓ **B. For (1), prompt for volume and divergence; for (2), prompt for sourcing and verification, not creative range**

Brainstorming benefits from breadth and divergent options; verifying a statistic benefits from precision and traceability to a source. Treating both tasks the same wastes the strengths of each approach.

Why not the others: *A ignores that the two tasks need opposite qualities. C inverts the correct pairing. D is false — task type meaningfully changes which prompting strategy performs better.*

1.4 — Correct: B

✓ **B. Context about the audience and purpose — that this is a board-level summary requiring a different level of detail than the source report**

An instruction without audience context forces Claude to guess at the appropriate level of detail. Naming the audience and purpose up front lets the first draft land close to usable.

Why not the others: *A was already satisfied. C is vague and does not specify the actual constraint (board-level, not merely shorter). D treats a prompting gap as a manual post-processing chore instead of fixing the prompt.*

1.5 — Correct: A

✓ **A. Provide two or three real examples of posts that capture the desired voice, alongside the request**

Concrete examples of the desired output are typically more reliable than an abstract description, especially for stylistically specific work like brand voice, because they show rather than describe the pattern to match.

Why not the others: *B has already been tried in a different form (a longer description) without success. C gives up on a solvable problem. D does not address the task at all.*

1.6 — Correct: B

✓ **B. The prompting effort is likely mismatched to the task's low stakes — a lighter, faster prompt would serve just as well**

Prompting effort should scale with what the output is used for. A low-visibility, low-stakes, easily corrected deliverable like internal meeting notes doesn't warrant the same investment as a high-stakes one.

Why not the others: *A and D both assume more effort is always better, which wastes time on a task where it doesn't pay off. C dismisses a reasonable use case outright.*

1.7 — Correct: B

✓ **B. Split the request into named sub-tasks — venue shortlist, draft agenda, vendor outreach email, budget tracker — and address each with its own prompt and review**

A 300-person conference has many independent work-streams. Naming them individually gives each one a clear deliverable and lets the coordinator review and adjust each piece before moving to the next.

Why not the others: *A repeats the same underspecified request. C removes the coordinator's judgment from a decision that should stay with them. D is an unrelated timing question, not a decomposition strategy.*

1.8 — Correct: A, C

✓ **A. Request the widest possible range of options for an open brainstorming task, rather than a single early favorite** / **C. For a drafting task, specify the intended audience and tone rather than leaving them unstated**

Divergent, high-volume output suits brainstorming; explicit audience and tone framing suits drafting. Matching the prompt's shape to the task type is the core skill this domain tests.

Why not the others: *B, D, and E all collapse a meaningful distinction between task types into a single one-size-fits-all approach, which is exactly what this domain warns against.*

Domain 2 — Output Evaluation and Validation (21.0%)

2.1 — Correct: B

✓ **B. Claude can generate fabricated content, including citations, in a tone indistinguishable from accurate content — formatting confidence is not evidence of existence**

A hallucinated citation can be perfectly formatted and read with total confidence. Formatting and fluency say nothing about whether the source actually exists — only checking against a real database does.

Why not the others: *A treats a surface property as a trust signal, which is the exact failure mode this domain warns against. C is false — hallucination is not limited to obscure topics. D misattributes a model behavior to a transcription error without evidence.*

2.2 — Correct: B

✓ **B. The statistic should be checked against an authoritative source before it is used externally**

A specific factual claim heading into an external presentation is exactly the kind of detail that needs verification against a real source — the stakes of being wrong publicly are high enough to warrant the check.

Why not the others: *A treats natural-sounding phrasing as sufficient verification. C disguises an unverified number rather than verifying it. D is an overcorrection — verifying is more useful than discarding a potentially correct fact.*

2.3 — Correct: B

✓ **B. A completeness failure — accurate content that silently omits material that was part of the original input**

Every stated sentence being accurate does not guarantee the summary is complete. Checking for what was left out is a distinct pass from checking whether what's present is correct.

Why not the others: A describes the opposite problem (fabricated addition, not omission). C ignores that an omission can materially mislead a reader even without a single false sentence. D mischaracterizes a substantive gap as a cosmetic one.

2.4 — Correct: B

✓ B. The summary may be biased in its selection and emphasis, even without containing any factual error

A summary can be composed entirely of accurate statements and still misrepresent the overall picture by over-weighting one segment of the input. Bias review is a check on emphasis and selection, separate from fact-checking.

Why not the others: A conflates factual accuracy with fair representation, which are different properties. C adds a specific fix without addressing the underlying selection bias. D misclassifies a substantive representativeness problem as cosmetic.

2.5 — Correct: B

✓ B. No — a regulatory response is high-stakes and externally visible, which calls for a second human reviewer regardless of how confident the draft reads

Escalation is driven by the stakes, reversibility, and visibility of the output — a regulatory response scores high on all three — not by how confident the draft feels to the person who read it.

Why not the others: A and C both substitute an irrelevant signal (confidence, length) for the actual escalation criteria. D overstates the rule; the issue is the review step, not a blanket ban on external AI-assisted drafts.

2.6 — Correct: B

✓ B. Not necessarily — accuracy and audience-fit are separate checks, and the technical version may need restructuring or simplifying for executives

A document can be fully accurate for its original audience and still be the wrong format, tone, or level of detail for a different one. Re-tuning for audience is a distinct step from verifying the facts.

Why not the others: A conflates factual accuracy with communicative fitness. C wrongly implies the facts themselves must change. D places an unreasonable burden on the audience instead of adapting the communication.

2.7 — Correct: A

✓ A. A structured table for the comparison; a quick inline chat answer for the one-off question

A structured table suits side-by-side comparison across many items; a quick, disposable question suits a plain inline answer. Matching format to how the result will actually be used is itself part of good output practice.

Why not the others: B over-engineers the one-off question. C under-serves the comparison task, which benefits from structure. D reverses the correct pairing entirely.

2.8 — Correct: B

✓ B. Compare both drafts against the same explicit criteria — accuracy, completeness, tone, and length — and choose based on that comparison

Comparing candidate outputs against the same stated criteria produces a decision that can be explained and defended, rather than one based on an impression that may not track what actually matters for the announcement.

Why not the others: A relies on an unreliable signal (fluency) rather than substantive criteria. C abdicates the decision rather than making it. D risks introducing inconsistency by merging without review.

2.9 — Correct: B

✓ B. The hedge was a genuine signal to verify that figure; removing it without checking discarded the very protection it was meant to provide

When Claude flags low confidence on a specific claim, that flag is useful information. Editing it away for tone, rather than resolving it by checking the fact, converts a known risk into a hidden one.

Why not the others: *A treats a cosmetic improvement as more valuable than resolving a real uncertainty. C only moves the same mistake earlier in the process. D dismisses a genuinely useful signal outright.*

2.10 — Correct: A, C

✓ **A. Confirming the memo addresses every criterion the original request specified, not only the ones easiest to answer / C. Reviewing whether the memo fairly represents all the vendors considered, rather than favoring the one the drafter expected to win**

Accuracy alone doesn't guarantee completeness (every requirement addressed) or balance (fair representation across options) — both are separate, necessary checks even once every stated fact has been confirmed correct.

Why not the others: *B is a surface-level check that doesn't address substance. D adds length without addressing either gap. E ignores that a validated output needs more than accuracy alone.*

2.11 — Correct: B

✓ **B. An internal consistency failure — the document contradicts itself even though each piece traces to a real source**

Internal consistency is its own check: a long document's figures, dates, and claims need to agree with each other from start to finish, independent of whether each one is individually traceable to a source.

Why not the others: *A misses that individually-correct figures can still contradict each other when combined incorrectly. C mischaracterizes a substantive contradiction as cosmetic. D misapplies the concept of bias to what is actually an arithmetic or aggregation error.*

2.12 — Correct: A, C

✓ **A. Whether it fairly represents both positive and negative feedback in proportion to what was actually collected / C. Whether the level of detail and tone suit a company-wide audience, or whether it was written more like an internal analyst note**

Correct quotation of figures does not guarantee balanced representation or audience fit — both require a distinct check even when the underlying numbers are accurate.

Why not the others: *B and D are surface-level or arbitrary constraints that don't address substantive validation. E wrongly treats accurate quotation as the entire validation task.*

2.13 — Correct: C

✓ **C. The stakes and reversibility of sending it to a customer in this specific situation — which call for review regardless of how the draft reads**

A customer situation involving potential legal action is high-stakes and hard to walk back once sent. That fact — not the drafter's confidence or the email's polish — is what should drive the decision to escalate for review.

Why not the others: *A and B both substitute a feeling or a surface impression for an assessment of actual stakes. D is an arbitrary and irrelevant proxy for risk.*

Domain 3 — Product and Model Selection (12.0%)

3.1 — Correct: B

✓ **B. A Project dedicated to the client, holding the reference material and standing instructions so they persist across conversations**

Recurring work with its own context, tone, and reference material is exactly what a Project is for: the setup persists so it doesn't need to be rebuilt in every new conversation.

Why not the others: A and C both recreate the same setup cost repeatedly instead of persisting it once. D is false — this is a standard, well-supported use case.

3.2 — Correct: B

✓ B. A research-oriented workflow suited to synthesizing multiple sources on an open-ended question

An open-ended question that benefits from drawing together multiple sources is a better fit for a research-oriented approach than a single quick chat turn, which is suited to narrower, more self-contained questions.

Why not the others: A undersells the open-ended nature of the question. C is not a strategy, just repetition. D artificially narrows a question that benefits from breadth.

3.3 — Correct: B

✓ B. Use an artifact so the proposal can be edited in place across rounds, rather than regenerated from scratch each time

A document that will be iterated on repeatedly is exactly what an artifact is designed for — it can be edited in place, keeping the working document consistent across rounds instead of being retyped each time.

Why not the others: A is inefficient and increases the chance of losing track of the latest version. C discards useful context between rounds. D unnecessarily rules out a well-supported use case.

3.4 — Correct: B

✓ B. A lighter, faster tier is likely the better fit, since the task is simple and high-volume, and using a heavier tier would add cost and latency without improving accuracy

A simple, well-defined, high-volume classification task is the textbook case for a lighter, faster tier: it meets the accuracy bar at a fraction of the cost and latency of the most capable tier.

Why not the others: A defaults to the most powerful option without matching it to the actual task, which wastes resources on this kind of task. C is factually wrong. D is not a coherent selection strategy.

3.5 — Correct: B

✓ B. Start a fresh conversation (or move to a Project with the relevant persistent knowledge) for the new, unrelated topic

Once a conversation has drifted across unrelated topics and grown long, starting fresh (or moving to a Project with the right persistent context) produces a more focused result than continuing to extend an already-cluttered conversation.

Why not the others: A keeps adding to a shared, finite resource already under strain. C is a less reliable substitute for simply starting a new conversation. D does not address the underlying clutter.

3.6 — Correct: B

✓ B. The actual binding constraint for this task — here, response speed — rather than a general preference for the most powerful option

The right choice follows from identifying which constraint actually matters for the task at hand. When latency is the binding constraint, optimizing for it — rather than defaulting to maximum capability — is the correct call.

Why not the others: A ignores the stated constraint entirely. C substitutes price for an actual assessment of fit. D dismisses a decision that has a clear right answer given the stated priority.

3.7 — Correct: B, C

✓ **B. Moving a recurring body of work with persistent context into a Project rather than repeating setup in every new chat / C. Matching model tier to the task's actual constraint — cost, speed, or quality — rather than to a fixed habit**

Persisting recurring context in a Project and matching model tier to the actual binding constraint are both core selection skills this domain tests.

Why not the others: *A wastes resources by ignoring tradeoffs. D lets a shared, finite resource degrade instead of managing it. E chooses the wrong container for iterative work.*

Domain 4 — Workflow Integration and Solution Design (16.0%)

4.1 — Correct: B

✓ **B. A human should confirm the requirements are complete and correctly prioritized before design work proceeds**

Claude can meaningfully accelerate organizing messy input into structured requirements, but confirming completeness and priority against what stakeholders actually meant remains a human responsibility the drafting step doesn't discharge.

Why not the others: *A skips a necessary human check. C overcorrects against a genuinely useful capability. D hands over a judgment call that should stay with accountable stakeholders.*

4.2 — Correct: B

✓ **B. The plan should still be checked against the same evaluation standards as any other output — accuracy, completeness, and fit for the actual situation**

Claude can substantially accelerate research and planning work, but the output still needs the same accuracy and completeness checks as any other Claude-assisted deliverable before it's acted on.

Why not the others: *A and C both substitute a surface impression (thoroughness, length) for actual validation. D is false — no task category is exempt from evaluation.*

4.3 — Correct: B

✓ **B. The reminder messages, because the process is describable and its inputs are consistent, unlike the triage decisions, which rest on undocumented tacit judgment**

A good process-optimization candidate has describable steps and consistent inputs. The reminder messages qualify; the triage decisions rest on undocumented, hard-to-transfer clinical judgment, which is a poor fit regardless of how valuable the decision is.

Why not the others: *A ignores a meaningful difference in fit. C picks the harder, less appropriate case based on value alone rather than suitability. D dismisses a genuinely good use case.*

4.4 — Correct: B

✓ **B. The process is likely still slow overall, since the acknowledgment doesn't resolve the unclear ownership causing the underlying delay**

Bolting a faster step onto a workflow broken by unclear ownership doesn't fix the workflow — it just produces a faster acknowledgment ahead of the same underlying delay. Effective integration means addressing the ownership problem itself.

Why not the others: *A and C both assume adding Claude automatically fixes a structural problem it didn't address. D is an unsupported and unrelated claim.*

4.5 — Correct: B

✓ **B. A thinking partner across iterations, proposing and revising a solution structure in response to feedback**

Proposing a structure and revising it against feedback across multiple rounds is exactly the kind of iterative solution-design support Claude can provide, distinct from producing a single unchangeable output.

Why not the others: *A mischaracterizes the interaction as one-shot when it's iterative. C is false given the scenario. D overstates Claude's role as replacing rather than supporting the team's judgment.*

4.6 — Correct: B

✓ **B. It oversells capability and omits the human oversight that should remain in place, which will erode trust once a limitation surfaces in production**

A stakeholder proposal should state plainly what a workflow change will and won't do, and what oversight remains. Overselling capability sets an expectation that fails visibly the first time a real limitation appears.

Why not the others: *A and C both encourage overselling, which backfires once reality falls short of the claim. D is false — communicating value and limitations accurately is explicitly part of this domain.*

4.7 — Correct: B

✓ **B. Run a small, reversible pilot at one or two stores first to surface failure modes before scaling further**

A small, reversible pilot on a real but bounded slice of the rollout reveals failure modes a paper plan doesn't, and keeps the cost of a wrong assumption low before it's multiplied across 40 locations.

Why not the others: *A and C both skip a low-cost step that would catch expensive mistakes before they scale. D avoids the risk entirely but also abandons the goal of a wider rollout.*

4.8 — Correct: B

✓ **B. An assigned owner responsible for monitoring the integration's output quality on an ongoing basis after the pilot ended**

A successful pilot with no assigned ongoing owner tends to degrade unnoticed as underlying conditions shift. Integration isn't finished at launch — someone needs to own its quality over time.

Why not the others: *A ignores that quality maintenance requires active ownership, not a one-time success. C addresses timing, not the actual missing element (ownership). D is contradicted by the scenario itself.*

4.9 — Correct: A, C

✓ **A. Piloting the change on a small, reversible slice of the process before scaling company-wide / C. Assigning an owner to monitor output quality on an ongoing basis after launch**

Piloting before scaling and assigning ongoing ownership are both core safeguards this domain tests — they catch failure modes early and prevent quiet quality decline after launch.

Why not the others: *B ignores the root cause of the workflow problem. D skips a low-cost risk-reduction step. E abandons the integration exactly when ongoing attention is most needed.*

4.10 — Correct: B

✓ **B. The underlying ambiguity is likely to persist or worsen, since adding a capability at one step doesn't redesign the surrounding process**

Inserting a capability into a step doesn't redesign the process around it. An unresolved ownership question upstream or downstream of that step will likely persist, since the integration didn't address it.

Why not the others: *A and C both assume technology resolves an organizational design problem on its own, which contradicts the domain's core distinction. D is false — this is precisely the kind of concern this domain addresses.*

Domain 5 — Configuration and Knowledge Management (12.0%)

5.1 — Correct: B

✓ B. In the Project's instructions, so it applies automatically to every conversation within the Project

Standing guidance meant to apply every time belongs in Project instructions, which persist and apply broadly, rather than being retyped by hand in every new conversation.

Why not the others: A requires manual repetition and is easy to forget. C is false — this is a standard, supported configuration option. D fails to actually inform Claude of the rule at all.

5.2 — Correct: B**✓ B. Upload the style guide and glossary as knowledge sources within the Project**

Uploading reference material as knowledge sources gives every conversation within the Project a stable, shared basis to draw on, instead of depending on what a person happens to paste in each time.

Why not the others: A is repetitive and error-prone compared to a one-time upload. C is false — internal, organization-specific terminology is not something Claude would already know. D does not configure Claude at all.

5.3 — Correct: B**✓ B. Use a connector to the shared drive so Claude references the live, current version whenever needed**

For material that changes frequently, a connector to the live source stays current automatically, which is more reliable than a document uploaded once and left to go stale.

Why not the others: A guarantees staleness as prices change. C is manual and error-prone. D risks relying on outdated information from an earlier session.

5.4 — Correct: B**✓ B. Instructions typed into a single message apply only to that message; a durable rule needs to be set at the system or Project level to persist across conversations**

A one-off message's instructions apply to that message, not automatically to future ones. A standing rule meant to hold going forward needs to be placed where it will persist — at the system or Project level.

Why not the others: A misunderstands how message-level instructions work; this is expected behavior, not a bug. C invents an unsupported explanation specific to competitors. D is not how persistence works.

5.5 — Correct: B**✓ B. Claude may reference the outdated version of the policy, since the Project's configuration is only as current as its last update**

Knowledge sources drift out of date as the underlying material changes. A Project's configuration needs an owner who updates it; a one-time upload does not stay current on its own.

Why not the others: A and D both incorrectly assume automatic freshness that isn't how uploaded documents work. C misunderstands how a knowledge source is used — it doesn't need to be manually opened to be referenced.

5.6 — Correct: B**✓ B. It becomes harder for Claude to apply the right material to a given task, and irrelevant content from one initiative can leak into an answer about another**

A Project should hold knowledge relevant to its own scope of work. Mixing unrelated bodies of knowledge into one Project increases the chance that irrelevant material influences an unrelated answer.

Why not the others: A and C both treat scope-mixing as beneficial when it actually degrades relevance. D mischaracterizes convenience as best practice.

5.7 — Correct: A, C

✓ **A. Placing standing, repeat-every-time guidance in Project instructions rather than a single chat message / C. Using a connector rather than a static upload for information that changes frequently**

Placing standing rules where they persist, and connecting to live sources for frequently-changing information, are the two configuration practices this domain rewards.

Why not the others: *B ignores the need for ongoing maintenance. D mixes unrelated scopes, which degrades relevance. E places a durable rule somewhere it won't persist.*

Domain 6 — Governance, Risk, and Responsible Use (15.0%)

6.1 — Correct: B

✓ **B. Drafting talking points is a reasonable use of assistance; the termination decision itself requires a specific accountable human and should not be delegated to Claude**

Being technically able to generate an answer says nothing about whether a use case is appropriate. A termination decision requires a specific accountable human by policy and by fairness; drafting preparatory talking points does not carry the same requirement.

Why not the others: *A ignores the meaningful difference between assistance and delegated accountability for a consequential decision. C dismisses a legitimate use case. D justifies an inappropriate delegation on the basis of convenience alone.*

6.2 — Correct: B

✓ **B. Health information carries the same handling obligations in a prompt that it carries anywhere else, and its use should be evaluated against those obligations before including it**

Data sensitivity rules don't relax because the processing step involves an AI system. Information that would need protecting anywhere else needs the same protection, and the same scrutiny, in a prompt.

Why not the others: *A assumes internal use removes an obligation that doesn't disappear that way. C is a false and risky distinction. D ignores that some data categories carry materially higher handling obligations than others.*

6.3 — Correct: B

✓ **B. The organization, which must know and apply the regulatory constraints on what can be shared with any AI system, including Claude**

Regulatory and privacy constraints are the organization's to know and apply to any system it adopts, including Claude — the technology does not enforce this on the organization's behalf.

Why not the others: *A misplaces responsibility onto the tool rather than the organization using it. C is factually incorrect. D isolates responsibility to an individual rather than the organization that sets policy.*

6.4 — Correct: B

✓ **B. Seek the required approval before proceeding, since the organization's own policy governs this decision regardless of technical capability**

An organization's own acceptable-use policy governs usage decisions even where the technology would allow more. Technical capability and organizational permission are separate questions, and policy compliance doesn't yield to convenience or urgency.

Why not the others: *A, C, and D all substitute a reason to skip the required approval for actually seeking it, none of which override a stated organizational policy.*

6.5 — Correct: B

✓ **B. No — the person who acted on and sent the output remains accountable for its correctness and consequences, regardless of what tool assisted in drafting it**

Using Claude to help draft something doesn't transfer accountability for the action taken on the result. The person who sent the figure to the client remains responsible for having verified it first.

Why not the others: A and C both attempt to shift accountability away from the person who took the action, which contradicts the principle that human accountability doesn't transfer to the tool. D makes accountability contingent on an unrelated factor.

6.6 — Correct: B

✓ **B. Not necessarily — an output can be fully compliant and still raise a fairness or ethical concern worth weighing on its own terms**

Compliance is the floor for responsible use, not the ceiling. A plan can clear every written rule and still raise a legitimate ethical concern — targeting a vulnerable population — that deserves separate consideration.

Why not the others: A and D both stop the analysis at compliance or performance, skipping the ethical dimension this domain explicitly covers. C incorrectly claims a technical limitation that isn't the actual issue.

6.7 — Correct: B

✓ **B. Whether to disclose that the content was AI-assisted, given that the audience would reasonably expect to know**

Where an audience would reasonably expect to know that content was AI-assisted, disclosure is a distinct consideration from whether the content itself is accurate — the two questions don't collapse into one.

Why not the others: A wrongly assumes accuracy alone settles the matter. C and D raise considerations unrelated to the transparency question actually at stake.

6.8 — Correct: A, C

✓ **A. Using Claude to draft a decision that, by policy, requires a specific accountable human to make / C. Including regulated personal data in a prompt without considering applicable data protection obligations**

Delegating a decision that requires specific human accountability, and including regulated data without considering the applicable obligations, are both governance concerns independent of output accuracy.

Why not the others: B, D, and E are all low-risk, appropriate use cases that don't raise the same governance concerns.

6.9 — Correct: B

✓ **B. Differently — drafting an internal explanation is a reasonable assistive use, while an individual claims-denial decision is a consequential, individually-impactful decision that calls for specific human accountability**

The same underlying tool can support very different requests with very different appropriateness. An internal explanatory draft is low-risk; an individual, consequential decision like a claims denial calls for the accountability and judgment of a specific human.

Why not the others: A collapses a meaningful distinction based only on industry and tool. C over-restricts a legitimate use case. D misunderstands the scope of regulatory concern, which extends well beyond marketing.

Domain 7 — Troubleshooting and Optimisation (10.0%)

7.1 — Correct: B

✓ **B. Diagnose the likely cause first — missing context, an ambiguous instruction, the wrong model tier, or a genuinely difficult task — before deciding what to change**

An off-target output can stem from several different causes, each with a different fix. Diagnosing which one is actually responsible comes before deciding whether to add context, clarify the instruction, change tier, or accept the task is genuinely hard.

Why not the others: A, C, and D all skip the diagnostic step and jump to an action that may not address the actual cause.

7.2 — Correct: B

✓ **B. It's now impossible to tell which of the four changes actually caused the improvement, and a new problem introduced by one of them could be hidden by the others**

Changing several variables at once may produce a better result, but it prevents attributing the improvement to a specific cause, and could mask a new issue introduced by one of the simultaneous changes.

Why not the others: *A treats the outcome as sufficient without considering whether the process is repeatable or diagnosable. C and D both incorrectly claim this approach is necessary or optimal, when isolating variables is the more reliable practice.*

7.3 — Correct: B

✓ **B. Add a standing instruction — at the system or Project level — specifying the desired length, so it applies automatically going forward**

A consistent pattern across many outputs points to a standing instruction that's missing, not a one-off correction to repeat by hand every time. Adding the rule where it will persist addresses the root cause.

Why not the others: *A treats a systemic pattern as a series of unrelated one-off fixes. C dismisses a fixable, recurring problem. D avoids addressing the actual cause.*

7.4 — Correct: B

✓ **B. The task may have outgrown the model tier in use, rather than the prompt itself being at fault**

When careful prompt refinement hasn't resolved a persistently weak result on a demanding task, that pattern can indicate a task-model mismatch — the task may call for a more capable tier, rather than a wording problem.

Why not the others: *A assumes the prompt is still the bottleneck despite repeated refinement attempts. C overstates the limitation as absolute rather than tier-dependent. D abandons a fixable situation prematurely.*

7.5 — Correct: B

✓ **B. The output should still be re-validated after the change, rather than assumed automatically equivalent to before**

Efficiency changes like shortening or reusing a prompt structure can shift the output in ways that aren't obvious until checked. Re-validating after the change, rather than assuming equivalence, keeps quality intact alongside the time savings.

Why not the others: *A wrongly assumes past reliability guarantees future reliability after a change. C is an overreaction without evidence of a problem. D removes the very check that would catch an issue.*

7.6 — Correct: A, C

✓ **A. Diagnosing the likely cause — context, instruction clarity, model tier, or task difficulty — before changing anything / C. Changing one variable at a time and re-checking the result before changing another**

Diagnosing the cause before acting, and isolating one variable at a time when adjusting, are the two disciplined troubleshooting practices this domain rewards.

Why not the others: *B reintroduces the variable-conflation problem this domain warns against. D gives up prematurely. E is not a troubleshooting strategy at all.*

Bonus — Extended Topics (15 answers)**B.1 — Correct: B**

(Prompting and Task Execution)

✓ **B. Draft the English version, then explicitly ask for a culturally adapted version for the Spanish-speaking audience, not a literal translation**

Adaptation for a different cultural audience is a distinct task from literal translation. Naming that distinction in the prompt — adaptation, not translation — produces a result that actually fits the target audience.

Why not the others: *A produces a technically translated but potentially culturally mismatched result. C discards useful continuity between the two versions. D ignores that cultural fit doesn't follow automatically from translation.*

B.2 — Correct: B

(Output Evaluation and Validation)

✓ **B. That the chart's values, axis labels, and scale accurately reflect the underlying data provided**

Visual polish says nothing about whether the underlying values are represented accurately. A chart needs the same accuracy check as any other output — its numbers, labels, and scale need to match the source data.

Why not the others: *A and D both substitute a surface impression for a substantive check. C addresses a cosmetic concern rather than the accuracy of what's being visualized.*

B.3 — Correct: B

(Product and Model Selection)

✓ **B. Use an editable, artifact-based or presentation-oriented workflow that supports revising the same deck across multiple sessions**

A deliverable that will be revised repeatedly over several days is best supported by a workflow that keeps the working version editable and persistent, rather than regenerating it from scratch or starting over each time.

Why not the others: *A and C both discard useful continuity between revisions. D dismisses a well-supported use case.*

B.4 — Correct: B

(Workflow Integration and Solution Design)

✓ **B. A shared Project (or shared knowledge sources) that both teams can draw on, keeping the underlying material consistent across their separate work**

When multiple teams need the same underlying reference material for related but separate work, sharing that material through a common configuration avoids the drift and duplication that comes from each team maintaining its own copy.

Why not the others: *A risks the two teams' copies diverging over time. C blocks a legitimate use case unnecessarily. D discards existing, presumably accurate material for no clear reason.*

B.5 — Correct: B

(Configuration and Knowledge Management)

✓ **B. A connector to the ticketing system, so current ticket status can be referenced as it changes**

Ticket status changes frequently, which is exactly the case where a connector to the live system is more reliable than a static, one-time export that will drift out of date almost immediately.

Why not the others: *A guarantees staleness on frequently-changing data. C is manual and unreliable at scale. D is false — this is a supported configuration pattern for frequently-changing sources.*

B.6 — Correct: B

(Governance, Risk, and Responsible Use)

✓ **B. Follow the publication's editorial policy and disclose the AI assistance as required**

Where an organization's own policy requires disclosure of AI involvement, that policy governs the decision — quality of the writing doesn't substitute for a required transparency step.

Why not the others: A and D both skip a disclosure the stated policy requires. C is a false absolute that ignores the specific policy in play here.

B.7 — Correct: B

(Troubleshooting and Optimisation)

✓ **B. Develop a reliable prompt template for the letter type, reuse it across instances, and periodically re-check outputs to confirm quality holds**

A proven, reusable prompt structure for a repetitive task saves real time, provided outputs are still periodically re-validated rather than assumed permanently correct once the template is established.

Why not the others: A wastes the efficiency gain available from a repetitive, well-defined task. C removes the ongoing check needed to catch drift. D adds length without addressing actual quality.

B.8 — Correct: B

(Output Evaluation and Validation)

✓ **B. Recognize this as scope creep — the output has grown beyond and diverged from what was actually asked — and revise it back toward the original request before using it**

An output that has quietly expanded beyond the original ask — especially by introducing new substantive claims — needs to be caught and corrected before use, regardless of how polished it reads.

Why not the others: A and C both accept unrequested, unverified additions at face value. D compounds the scope creep instead of correcting it.

B.9 — Correct: B

(Governance, Risk, and Responsible Use)

✓ **B. This breaches a contractual disclosure obligation, which is separate from whether the report's content is accurate**

A contractual disclosure requirement is a distinct obligation from content accuracy. Meeting one doesn't satisfy the other, and an accurate report can still breach a disclosure term the client is owed.

Why not the others: A and D both wrongly treat accuracy as sufficient regardless of a separate contractual obligation. C invents an exception that doesn't follow from the scenario.

B.10 — Correct: B

(Prompting and Task Execution)

✓ **B. Compare the Claude-drafted response against a sample of human-written responses using the same explicit criteria (accuracy, tone, resolution rate)**

A defensible comparison uses the same explicit criteria across both sources, rather than assuming one is better by default or judging on an arbitrary proxy like length or speed of production.

Why not the others: A and D both substitute an irrelevant signal for a substantive comparison. C wrongly claims this kind of benchmarking isn't possible.

B.11 — Correct: B

(Workflow Integration and Solution Design)

✓ **B. Supporting the human decision-maker by organizing the facts and surfacing considerations, while the final decision itself stays with the senior human as policy requires**

Recognizing when a task calls for a specific accountable human doesn't mean Claude has no role at all — it means the decision itself stays with that person, while assistance with organizing information remains appropriate.

Why not the others: A and D both let the technology make or effectively make a decision policy reserves for a human. C over-restricts a genuinely useful supporting role.

B.12 — Correct: B

(Product and Model Selection)

✓ **B. Summarize or reset the conversation's active context to the current phase, while keeping the Project's durable knowledge sources and instructions intact**

A cluttered conversation can be refreshed without discarding the Project's durable configuration — summarizing or resetting the active context addresses the clutter while preserving the setup that still applies.

Why not the others: A lets an increasingly irrelevant history keep influencing new answers. C discards durable, still-useful configuration along with the clutter. D abandons a Project that's still relevant to the ongoing initiative.

B.13 — Correct: B

(Configuration and Knowledge Management)

✓ **B. Sharing at the Project level can expose sensitive material to people who shouldn't have access to it, so scope and access need to be considered together, not just convenience**

Configuring a Project's scope and knowledge sources includes considering who has access to what's shared — broad convenience-based sharing can expose sensitive material to people who shouldn't see it.

Why not the others: A ignores a real access problem. C overcorrects into a blanket rule not supported by the scenario. D artificially separates two considerations (access and sensitivity) that are linked here.

B.14 — Correct: A, C

(Governance, Risk, and Responsible Use)

✓ **A. Whether customers interacting with the bot would reasonably expect to know they're talking with an AI system / C. Whether the bot might give financial or legal-sounding advice it isn't qualified to give, without a human escalation path**

Transparency about AI involvement and avoiding ungoverned advice-giving without an escalation path are both responsible-use considerations, independent of the bot's technical implementation or appearance.

Why not the others: B, D, and E are implementation or performance details that don't bear on responsible use in the way transparency and escalation do.

B.15 — Correct: A, C

(Troubleshooting and Optimisation)

✓ **A. The task may call for a more capable model tier than the one in use / C. The task itself may be a poor fit for AI assistance in its current form, and may need to be restructured or handled differently**

After careful, disciplined single-variable troubleshooting has failed to resolve a persistent problem, a genuine task-model mismatch or a poor task fit for AI assistance become reasonable explanations to consider.

Why not the others: B contradicts sound single-variable troubleshooting practice. D and E both stop investigating despite a clearly unresolved, persistent problem.